CSI-160 Python Programming
Foundations of Computing

Vikas Thammanna Gowda

07/08/2025

Contents

1 What Is a Computer?

1.1 Basic Hardware vs. Software

A **computer** is essentially a machine that can be *programmed* to carry out sequences of instructions and calculations automatically. Every computer system is built from a combination of **hardware** and **software**. *Hardware* refers to the physical components of a computer that you can touch – such as the CPU, memory chips, hard drive, keyboard, or monitor. In contrast, *software* refers to the programs and instructions that run on the hardware and tell the computer what to do. For example, an application like a web browser or a game is software, whereas the processor and memory it uses are hardware.

1.2 How CPU, Memory, Storage, and I/O Work Together

A computer's hardware components work in tandem to execute software instructions. The **Central Processing Unit (CPU)** – often called the "brain" of the computer – fetches and executes instructions from memory. The **memory** (RAM) holds data and program instructions that are actively being used, providing the CPU quick access to this information. The CPU and memory are connected via the motherboard and communicate over buses (electronic pathways). When you run a program, the program's instructions are loaded from long-term **storage** (like a hard disk or SSD) into memory, and the CPU reads those instructions from memory one by one to perform them. **Input/Output (I/O)** devices allow the computer to interact with the outside world: *input devices* (keyboard, mouse, etc.) let users or other systems send data into the computer, while *output devices* (monitor, printer, speakers) display or transmit the results of the computer's computations to users. All these parts are coordinated by the system's architecture to work together – for instance, when you press a key (input), the CPU processes that input, possibly stores or retrieves data from memory or storage, then outputs a result to your screen (output).

Figure 1: Block Diagram of a Computer

1.3 Types of Computers

Computers are often categorized by the scale of tasks they're designed to perform and the environments in which they operate. At one end of the spectrum are Personal Computers, built for individual users handling everyday tasks such as web browsing, document editing, and media playback. In the middle lie Workstations, which pack extra processing muscle and specialized hardware for demanding technical or creative work. At the top are Mainframes, colossal systems engineered to support thousands of users and process vast volumes of critical data with near-perfect reliability.

1.3.1 Classification Criteria

When deciding whether a machine is a PC, a workstation, or a mainframe, engineers look at several factors:

- **Processing Capability:** Personal Computers use consumer-grade CPUs and possibly a modest GPU. Workstations step up to server-class processors (often with many cores or specialized accelerators) and professional GPUs. Mainframes employ proprietary multi-chip architectures that deliver massive parallelism.

- **Memory and Storage:** A typical PC has a few gigabytes of RAM and one or two terabytes of disk or SSD. Workstations raise that to tens or hundreds of gigabytes of error-correcting (ECC) memory and high-speed storage arrays. Mainframes boast terabytes of memory and petabytes of redundant, enterprise-grade storage.

- **Reliability and Uptime:** Desktop PCs are built for an average of a few years of consumer-level use. Workstations add features like ECC memory, RAID disk arrays, and robust cooling to reduce downtime. Mainframes target "five nines" availability (99.999

- **User Capacity:** One person works at a PC at a time. A workstation typically serves a single specialized user but can sometimes share resources. A mainframe simultaneously supports hundreds or thousands of users running business-critical applications.

- **Cost and Support:** PCs are the most affordable category, ranging from a few hundred to a couple thousand dollars. Workstations cost several thousand to tens of thousands, reflecting their enhanced hardware and warranties. Mainframes carry six-figure price tags plus ongoing maintenance contracts.

1.3.2 Personal Computers

Personal Computers are optimized for individual productivity. A desktop PC—the classic tower or small-form-factor box—lets users upgrade components like RAM, storage, or graphics cards. A laptop integrates keyboard, display, battery, and trackpad into one portable package; its design balances performance against battery life and weight. The all-in-one PC (for example, Apple's iMac) houses all system components behind the display, offering a sleek footprint at the expense of internal expandability. Finally, tablets or convertibles run mobile or desktop operating systems on touchscreen hardware, trading raw power for extreme portability and instant-on convenience.

1.3.3 Workstations

Workstations sit between PCs and mainframes in both power and price. An engineering workstation typically includes multi-core server CPUs (such as Intel Xeon or AMD Threadripper), large banks of ECC-protected RAM, and high-performance storage—ideal for tasks like finite-element analysis or 3D CAD modeling. A graphics or video workstation pairs professional-grade GPUs (e.g., NVIDIA Quadro or AMD Radeon Pro) with NVMe SSD arrays, enabling real-time rendering, animation, and video editing workflows. In the era of data-driven science, data science/AI workstations bring together powerful CPUs, multiple GPUs connected by high-bandwidth interconnects, and ample memory so you can train machine-learning models right at your desk.

1.3.4 Mainframes

At the enterprise level, mainframes power the backbones of banking, insurance, government, and large-scale commerce. Transaction-processing mainframes handle enormous volumes of small, quick operations—such as ATM withdrawals or online retail purchases—without missing a beat. Batch-processing mainframes run scheduled jobs like payroll calculations or billing cycles, where hundreds of millions of records are processed in a single, windowed batch run. For extreme continuity, high-availability mainframes use redundant hardware clusters and advanced failover software so that even if one component fails, the system as a whole keeps running with virtually zero downtime.

2 What Is Programming?

Programming is the process of giving a computer precise instructions to perform tasks. Because computers cannot think for themselves, a programmer must write out step-by-step commands in a form the machine can follow exactly. In essence, programming involves designing a solution to a problem and then implementing that solution in a **program** – a set of instructions executed by the computer. We often start by formulating an **algorithm**: a sequence of steps or a procedure to solve a specific problem or accomplish a task. Algorithm consists of three parts:

- **Input(s):** The data or resources the algorithm requires.

- **Output(s):** The result(s) the algorithm produces.

- **Steps:** A precise, numbered list of actions that transform the inputs into the outputs.

By defining an algorithm in this way, we eliminate ambiguity and ensure that both humans and computers can follow the same procedure. Once an algorithm is defined, programming means translating that algorithm into a program using a programming language so that the computer can execute it. Good programming also requires being very exact – computers will do exactly what they are told, no more and no less. This means a programmer needs to think through every detail of the task. The process from conceptual *algorithm* to actual *program* involves careful planning (sometimes using pseudocode or flowcharts to sketch out the logic) and then coding, testing, and refining the instructions until the computer behaves as expected.

Algorithm: MakeTea

Input(s):

- Cold or room-temperature water, Tea leaves (or tea bag), Kettle or pot, Cup, (Optional) Sugar and/or milk

Output(s):

- A hot cup of tea, ready to drink

Steps:

1. Pour water into the kettle or pot.

2. Heat the water until it reaches a rolling boil.

3. Place the tea leaves or tea bag into the pot (or directly into the cup).

4. Pour the boiling water over the tea.

5. Let the tea steep for the desired time (typically 3–5 minutes).

6. Remove the tea leaves or tea bag from the liquid.

7. (Optional) Add sugar and/or milk, then stir until dissolved.

8. Serve the tea in a cup.

---

3 Programming Languages: Low-Level vs. High-Level

Not all programming languages are the same; they exist on a spectrum from low-level to high-level. The difference lies in how close the language is to the computer's hardware instructions versus how close it is to human-friendly abstractions.

3.1 Low-Level Languages (Machine Code and Assembly)

Low-level languages are languages that are very close to the hardware's native language (machine code). *Machine code* is the most fundamental level of code – binary digits (0s and 1s) that directly control the CPU's operations. This is extremely difficult for humans to read or write. *Assembly language* is a small step above machine code: it uses short textual mnemonics (like `ADD` for addition, `MOV` for moving data) instead of raw binary, but each assembly instruction still corresponds very closely to a single machine instruction. Low-level code gives the programmer very fine-grained control over the hardware (for example, managing memory addresses and CPU registers directly). However, it's not intuitive for humans – even simple tasks require many detailed instructions. For instance, printing the word "HELLO" in pure binary machine code would require a sequence of binary values representing each operation and character, something like `01001000 01000101 01001100 01001100 01001111` (which is not at all obvious to a human). In assembly, it would be somewhat more readable but still quite complex, involving multiple instructions to output each character. Low-level languages are generally used when maximum performance or direct hardware control is needed (such as writing device drivers or embedded systems software).

3.2 High-Level Languages (Python, Java, C++)

High-level languages, on the other hand, are designed to be closer to human language and abstract away much of the hardware detail. Languages like Python, Java, and C++ allow a programmer to accomplish tasks with more succinct and readable code. For example, to display the word "HELLO" on the screen in a high-level language, you might just write a single statement like `print("HELLO")` in Python, which is far easier to understand and write than the equivalent low-level code. High-level languages manage many aspects automatically – things like memory management, complex CPU instructions, etc., are handled behind the scenes. This makes developers more productive and programs more portable, as the same high-level code can often run on different types of machines with minimal or no modification. The trade-off is that high-level languages may be less efficient in terms of performance and may not allow as much fine control over the hardware. In practice, however, high-level languages are by far the most commonly used for application development because they greatly speed up development and reduce the chance of errors.

3.3 Trade-offs: Control vs. Productivity

The choice between low-level and high-level languages involves a trade-off between control and productivity. Low-level languages offer **greater control** over hardware and potential for highly optimized performance, but writing code in them is labor-intensive and error-prone. High-level languages prioritize **developer productivity** and ease of use, at the cost of some control and efficiency. In other words, it typically takes much less time and code to develop a program in a high-level language, and that code will be easier to read and maintain, whereas a low-level language might yield a faster or more optimized program but require significantly more effort to write. Modern software development often finds a balance: critical low-level components (where performance or hardware access is crucial) might be written in C or assembly, while higher-level application logic is written in languages like Python or Java. Understanding this spectrum helps programmers choose the right tool for a given task.

4 Program Translation

When we write a program in a high-level language, the computer cannot directly understand it. The code must be translated into machine instructions that the hardware can execute. There are several ways this translation happens, and it's important to know the distinctions:

4.1 Assemblers (for Assembly Language)

An **assembler** is the simplest translator, used for low-level assembly language. Assembly language uses human-readable mnemonics but is specific to a CPU's instruction set. The assembler program takes each assembly instruction and converts it into the corresponding machine code (binary instruction) for that CPU. Essentially, it has a one-to-one (or very close) mapping from assembly to machine code. The result of assembly is an executable machine code program. For example, if an assembly code has a line to add two numbers, the assembler will translate that into the exact binary opcode that tells the CPU to perform addition. Assemblers produce very efficient code and give access to hardware, but the programmer must manage every detail.

4.2 Compilers (e.g., for C/C++) vs. Interpreters (e.g., for Python)

A **compiler** is a program that translates code written in a high-level language into machine code (or an intermediate low-level form) *before* the program is run. It takes the entire source code, processes it, and produces an executable file or binary. For example, C and C++ are traditionally compiled languages: you write the code, run a compiler, and it outputs an executable program. This machine code can then be executed directly by the computer's CPU. The compilation process typically involves multiple stages (parsing the code, optimizing, etc.), but the key point is that after compilation, you have a standalone machine-code program. Compiled programs generally run very fast because the translation to machine instructions is done upfront.

An **interpreter**, by contrast, executes a high-level program by reading it one piece at a time and performing the operations on the fly, without a separate compile step. In an interpreted language (like standard Python, Ruby, or JavaScript), you feed the source code to an interpreter, and it handles each instruction step-by-step: read a line or statement, translate it (for example, into some internal form), and execute it immediately, then move on to the next. This means you don't get a separate binary file; the translation and execution happen together. Interpreted execution can be more flexible (you can often run code interactively, and it's easy to test changes quickly), but it tends to be slower than running compiled code because translation happens as the program runs, and some optimizations may not be as aggressive.

To illustrate, consider a simple program that calculates a result. In a compiled language like C, you would compile the program once (catching many errors at compile time), and then you can run the resulting executable as many times as needed efficiently. In an interpreted language like Python, you would run the Python interpreter on your script each time; the interpreter checks each line and executes it, which adds overhead during execution. Many scripting and beginner-friendly languages use interpreters to allow rapid development and debugging (since you can run code immediately without a compilation step).

4.3 Bytecode and Just-In-Time Compilation (JIT)

Modern language implementations often blend these approaches for efficiency. Some languages compile source code into an intermediate form known as **bytecode**. Bytecode is a lower-level, platform-independent code that is not as detailed as machine code but is easier to translate than high-level code. For example, Java is compiled into bytecode (.class files) which run on the Java Virtual Machine (JVM). Similarly, Python (when using the standard CPython interpreter) compiles source files into bytecode (with `.pyc` files) which its virtual machine then interprets. Bytecode improves performance over interpreting raw source because the high-level parsing is done ahead of time, and it also allows portability (the same bytecode can run on any machine that has the virtual machine for that language).

To further improve performance, many runtime systems use **Just-In-Time (JIT) compilation**. JIT compilers take the bytecode (or other intermediate representation) and compile parts of it into machine code *at runtime*, i.e., while the program is running. The idea is to get the best of both worlds: the flexibility of interpretation and the speed of compilation. For instance, the Java JVM will interpret the bytecode at first, but as it identifies "hot" code paths that run frequently, it will JIT-compile those into native machine code for faster execution. This means that after a short warm-up period, a Java program can run nearly as fast as a fully compiled program, because its most critical parts have been turned into optimized machine code by the JIT. Python itself has alternative implementations (like PyPy) that use JIT compilation to accelerate execution. In summary, program translation might involve pure ahead-of-time compilation, pure interpretation, or hybrid approaches (bytecode interpretation and JIT), each with advantages in terms of speed, portability, and development ease.

5 The Software Development Life Cycle

Writing a program (or developing any software) is not just about writing code. It involves a series of stages collectively known as the **Software Development Life Cycle (SDLC)**. The SDLC provides a structured approach to building software, ensuring quality and manageability throughout the process. We will outline the key phases of SDLC and what happens in each:

5.1 Requirements and Design

The first step is determining **what the software needs to do** and **how it should be designed to do it**. In the *requirements* phase, developers (often along with stakeholders or clients) gather and document the specifications: what problem is the software solving? What are the features and constraints? This often results in a Software Requirements Specification (SRS) document or a clear list of requirements. Once the team understands the goals, the next step is *design*. In the design phase, developers plan the structure of the software and how it will meet the requirements. This can include creating high-level architecture diagrams, flowcharts, and using pseudocode to sketch out algorithms before actual coding. For example, if you were designing a simple program like a calculator app, the design might include a flowchart of how user input flows through the system and a description of modules or functions (addition module, display module, etc.). The guiding principle in this stage is to figure out the blueprint of the software – breaking the problem into smaller components, deciding how those components will interact, and ensuring that all requirements will be addressed by the design. Good design often uses techniques like modularization (which we discuss later) to keep the system organized. By the end of this stage, the team should have a clear plan (perhaps documented as specifications, pseudocode, or diagrams) for how to build the software.

5.2 Implementation (Coding)

In the implementation phase, the developers actually **write the code** according to the design. This is the programming part: using a programming language to create the functions, classes, and modules that were planned in the design stage. The goal is to translate the design into a working software product. Developers will follow best practices for coding and use the chosen technology stack to build the application. During implementation, it's common to do incremental coding and frequent testing of small parts to ensure each part works correctly as it's built. The result of this phase is a working software system (or at least individual components of it) built according to the design. If we use the calculator app example, this is where developers write the actual code for the add, subtract, multiply, divide functionalities, and the user interface code, etc., following the plan made earlier. It's important that the code is written with care for *readability* and *maintainability* so that future developers (or the same developers later) can understand and modify it. The output of this phase is typically one or more program files (source code) that implement the required features.

5.3 Testing and Debugging

Once some code is written, it must be **tested** to ensure it works as intended. In the testing phase, developers (and possibly dedicated QA testers) run the program with various inputs and scenarios to verify that each requirement is met and to catch any **bugs** (errors or unexpected behavior). Testing can be broken down into multiple levels:

- *Unit testing*: checking individual components or functions for correctness.

- *Integration testing*: checking that different parts of the system work together properly.

- *System testing*: checking the entire integrated application against the requirements.

- *User acceptance testing*: sometimes, having actual end-users or clients test to see if the system meets their needs.

When tests reveal a problem, the process of **debugging** begins – this means locating the source of the error in the code and fixing it. Debugging can be challenging, as it requires understanding what the code is actually doing versus what it was intended to do. Tools like debuggers or simply strategic printouts/logs are used to inspect the program's state and find where it goes wrong. This stage is critical: it's much cheaper and easier to fix problems before software is deployed to real users. A well-tested and debugged program will be more reliable and maintainable. Testing and debugging continue iteratively until the developers are confident that the software is correct and stable.

5.4   Deployment and Maintenance

After testing is successful and the software is deemed ready, it is **deployed** – delivered to the end users or moved into a production environment. Deployment might mean publishing a downloadable application, releasing an update through an app store, or installing a system on a server for users to access. Sometimes deployment is done in stages (for example, a beta release to a small group before a full launch).

Once users begin using the software, the project enters the **maintenance** phase. Maintenance involves updating the software to fix any new bugs that appear in real-world use, improving features, or adapting the software to new requirements or environments over time. No software is ever truly "finished" – operating systems update, user needs evolve, and new security vulnerabilities might be discovered, so developers often need to issue patches or new versions. Maintenance also covers activities like optimizing performance after seeing how the system behaves in production or refactoring code for better clarity and extensibility. Essentially, maintenance ensures the software continues to meet user needs and run smoothly after its initial release. This phase can last as long as the software is in use.

It's worth noting that there are various models such as Waterfall, Agile, etc., that organize these stages differently – some are strictly sequential, while others repeat these stages in cycles – but nearly all software development will include these fundamental activities of planning, coding, testing, and maintenance in some form.

6 Basic Program Design Principles

To write good software (not just code that "works," but code that is robust, readable, and maintainable), programmers follow several important design principles. Here we discuss a few fundamental ones: **modularity**, **readability**, and practices like **incremental development** and **version control**.

6.1 Modularity: Functions and Modules

Modularity is the design principle of breaking a program into separate parts (modules) such that each part handles a specific piece of the overall functionality. In practice, this means using functions, classes, and modules (files or libraries) to encapsulate different functionality. For example, if you are writing a game, you might have one module or section of code for the player input, another for the game physics, another for rendering graphics, etc. Each module can be developed and understood independently, as long as the modules have well-defined interfaces (ways they interact with each other). This separation of concerns makes programs easier to understand and maintain because you can focus on one part at a time and because changes in one module (as long as the interface remains the same) won't heavily impact other parts. Functions are a basic unit of modularity: each function performs a well-defined task, and by calling functions, you can avoid repeating code and make the program structure clearer. In Python, for example, you can group related functions into a module (a .py file) and reuse that module in different programs. Overall, modular design leads to code that is more organized, reusable, and testable.

6.2 Readability: Naming, Comments, Style Guidelines

Code is read much more often than it is written. **Readability** refers to how easily a human (including the original programmer, weeks or months later) can understand the code. This is vital, because code that is hard to read tends to breed bugs and is difficult to extend or fix. In fact, programmers spend a majority of their time reading and understanding existing code rather than writing new code. Unreadable or sloppy code can lead to mistakes, inefficiencies, and even duplicate code when future developers can't tell that a feature already exists and end up reimplementing it.

Key practices to enhance readability include:

- **Meaningful naming**: Use clear, descriptive names for variables, functions, and other identifiers. For instance, a variable named `total_price` is more informative than `x`. Good naming makes the code self-documenting.

- **Comments and documentation**: Provide comments to explain non-obvious parts of the code or the rationale behind certain decisions. For example, a brief comment above a complex block of code can help others (or your future self) understand its purpose. However, comments should supplement clear code – they are not a substitute for it.

- **Consistent style and formatting**: Following a style guide (like PEP 8 for Python) ensures consistency in things like indentation, bracket placement, spacing, etc. Consistent indentation and formatting help visually structure code blocks and make it easier to parse code at a glance.

- **Modular decomposition**: As mentioned, breaking code into functions and modules not only helps design but also readability. Each function should ideally do one thing, and do it well (the Single Responsibility Principle), which makes it easy to understand.

In essence, writing readable code means writing code for *other humans* to understand, not just for the computer to execute. Following naming conventions, writing clear comments, and adhering to coding style guidelines are all practices that contribute to code quality and maintainability.

6.3 Incremental Development and Version Control

Developing software is an iterative process. **Incremental development** is a practice where you build and test your program in small pieces, gradually adding functionality, rather than trying to write the

entire program in one go. This approach involves writing a small amount of code, then running it to test and verify it works, then continuing to add more. By proceeding in increments, you catch errors early and avoid being overwhelmed by large amounts of new code that haven't been tested. For example, if you're creating a simple application, you might first implement the core function with a simple output and test it. Once that works, you add the next feature, test again, and so forth. This approach not only makes debugging easier (since you know any new problems are likely in the code you just added), but also gives you a working program (albeit with limited features) at each step. It aligns with the idea of "building up" a program iteratively and is especially useful for beginners to gain confidence and for large projects where continuous feedback is valuable.

Another essential practice in modern software development is using a **version control system** (VCS) such as Git. Version control systems help manage changes to the source code over time and facilitate collaboration between multiple developers. With version control, every change you make to the code is tracked. You can commit changes with a message describing what was done, and the system keeps a history of all these commits. This means you can revert to a previous version if a new change introduces a bug, or compare different versions of code to find out when a bug was introduced. For team projects, VCS allows multiple people to work on different parts of the code simultaneously without overriding each other's work – changes can be merged together systematically. It also provides a level of backup: the code repository exists on developers' machines and often on a remote server, so the risk of losing code is minimized. In short, version control is a safety net and an organizational tool; it enforces discipline in saving work and documenting changes, and is indispensable in professional software development. Even for an individual working solo, using version control is considered a good practice to keep their project history and progress well-managed.

7   Why Python?

Python is often recommended as an excellent language for beginners and experts alike, and it's worth understanding **why** – especially in the context of learning programming fundamentals. Python is a high-level programming language, and it exemplifies many of the concepts we've discussed in terms of ease-of-use and abstraction.

7.1   Simplicity, Readability, and Versatility

Python is known for its **simple and readable syntax**. This means that Python code often looks closer to pseudocode or even plain English, making it easier to learn and understand. For example, to get the length of a list in Python, you might write `len(my_list)`, and to print something, you just write `print("something")`. There are minimal unnecessary symbols or complicated structures for a beginner to worry about. Python enforces indentation for blocks of code (like loops or function definitions), which naturally leads to clean and consistently formatted code. These features lower the barrier to entry for new programmers and allow you to focus on learning programming concepts rather than complex syntax. In fact, Python's design philosophy (as described in "The Zen of Python") emphasizes readability and simplicity.

Despite its simplicity, Python is a very **powerful and versatile** language. It comes with a large standard library and has a vast ecosystem of third-party libraries (via the Python Package Index) for almost any task: web development, data analysis, artificial intelligence, automation, scientific computing, and more. This means that once you learn the Python basics, you can apply it to many different domains without having to learn a new language for each domain. For instance, a single Python program could grab data from the web, crunch numbers using a math library, and then output a graph – all using libraries readily available to Python. Because of this versatility, Python is used by companies and researchers in many fields, and Python skills are highly in demand.

Another reason Python is a great learning language is its **community and support**. There is extensive documentation, and a large community of Python users means that if you encounter a problem, it's likely someone else has as well – and you can find help on forums like Stack Overflow or through countless tutorials. Python's community also contributes to its many libraries and frameworks, constantly improving the tools available to programmers.

7.2   Python's Place in the Low-Level vs. High-Level Spectrum

As discussed earlier, Python is a **high-level language**. It abstracts away most of the low-level details of the computer. For example, Python manages memory for you (via garbage collection), it doesn't require you to manually handle addresses or allocate memory as you might in C. It allows you to write code without thinking about the specifics of the hardware. This makes Python highly productive – you can often implement a feature in Python in a fraction of the number of lines it would take in a lower-level language. The trade-off is performance: Python code is generally slower to run than equivalent code in a low-level language like C, because Python code is executed by an interpreter and does a lot of work behind the scenes. However, for many applications, especially at the learning stage, this is not a problem. Computers are fast, and Python is "fast enough" for tasks like scripting, web development, or prototyping. Moreover, when performance is critical, Python allows integration with low-level languages (for example, heavy computations can be offloaded to libraries written in C/C++ which Python can call).

In terms of spectrum, Python sits at the high-level end, emphasizing ease of use over fine-grained control. It means that as a learner, you can focus on *concepts* like loops, functions, and data structures without getting bogged down by complex syntax or machine-level details. Python also exemplifies many good programming practices – its emphasis on readability and the enforced indentation teach you to write clean code from the start.

To summarize, we learn Python not only because it's beginner-friendly (simple and readable) but also because it's a language with real-world strength: it's used in industry and academia, and it allows a gentle introduction to programming concepts that scale up to serious applications. Python serves as a gentle stepping stone into the wider world of programming: once you understand the concepts using

Python, you can transfer those concepts to other languages (the differences will mostly be in syntax and level of manual control). Meanwhile, you can also accomplish a lot directly with Python, making it a practical choice for projects big and small.